

Mapping AI in Scholarly Communication: Research Workflows, Named Models, and Governance

Jedson Gabriel Ferreira de Paula, Joaquim de Oliveira Vieira,
Willian Francisco da Silva, Claudio R. M. Mauricio
Universidade Estadual do Oeste do Paraná (UNIOESTE)
Foz do Iguaçu, Paraná, Brazil

Abstract—AI is reshaping scholarly writing, peer review, evidence synthesis, and institutional governance, but bibliometric mapping of these practices is frustrated by broad database searches that return thousands of clinical, engineering, and educational records. From 6,261 multi-source candidate records in a locally preserved April 2026 snapshot, we applied documented scope rules and retained 711 studies with publication years 2020–2026. Studies published in 2026 provide provisional context. We positioned the retained studies on two theory-guided semantic axes using prototype-anchored BGE-M3 embeddings. Technological specificity (T) runs from generic AI to named generative models; governance orientation (G) runs from workflow support to integrity and institutional policy. A greedy variant-filtering procedure selected prototypes to minimise the axis cosine (final cosine = +0.0007) with leave-one-out stability $\rho \geq 0.98$. T and G were positively associated after controlling for publication year (unstandardised $\beta = 0.206$, 95% CI [0.119, 0.300]; partial $r = 0.165$, $p < 0.001$). Within a blinded, score-stratified two-rater validation sample ($N = 129$), unweighted Spearman correlations between human ratings and final automated scores were 0.755 (T) and 0.623 (G). Named-model specificity was weakly positively associated with governance orientation, while workflow-support research remained present across the technological-specificity range.

Keywords—Bibliometrics; scholarly communication; generative artificial intelligence; research integrity; semantic embeddings; evidence synthesis.

I. INTRODUCTION

Generative AI has become a practical and debated component of academic work. Tools based on large language models can assist with drafting, searching, screening, summarising, and reviewing; the same tools have heightened debates over fabricated references, disclosure, authorship, assessment, and research integrity. These developments concern Latin.Science’s interests in artificial intelligence, governance, open technologies, and social impacts of technology.

The literature is difficult to characterise with a conventional keyword count. Terms such as *review*, *screening*, *writing*, and *AI* also occur in clinical AI, technical applications, and language-learning studies. A broad retrieval therefore

risks measuring the database’s ambiguity rather than the scholarly-communication field. Bibliometric work requires a transparent definition of the population, documented data-processing procedures, and claims matched to the information that metadata and citations can support [1]–[3].

This study addresses that problem by combining a documented scope-screening procedure with two interpretable semantic measurements. Rather than treating clusters from an embedding space as discovered topics, we define two dimensions that correspond to substantive questions. The first captures whether a paper concerns generic AI or named generative models. The second distinguishes work centred on task support from work centred on governance, integrity, and institutional oversight.

We map a screened, multi-source corpus with publication years 2020–2026. Five prototype abstracts per pole define each continuous axis. Variant selection minimises prototype-level overlap while preserving leave-one-out stability. We describe publication cohorts and estimate the association between technological specificity and governance orientation.

We address three research questions:

- 1) What publication-volume and semantic-composition patterns are observed across publication cohorts from 2020 to 2026 in this corpus snapshot?
- 2) How are generic/named AI and workflow/governance orientations distributed in the retained literature?
- 3) Is technological specificity associated with governance orientation after accounting for publication year?

II. RELATED WORK

Bibliometrics quantitatively studies bibliographic material to describe the structure and evolution of a research field [1]. Contemporary workflows combine transparent retrieval, deduplication, descriptive indicators, and visual or network-based analyses [2], [3]. The present study uses reporting principles from PRISMA 2020 and PRISMA-S as transparency

guides for documenting identification and selection; it is not presented as a systematic review or meta-analysis [4], [5].

Recent work on AI in scholarly communication has examined the rules and practices surrounding authorship, disclosure, and editorial work. In a study of 300 leading journals, Lund and Naheem found that more than half had adopted AI-authorship policies, but that the requirements differed by publisher and discipline [8]. The ICMJE recommendations similarly place responsibility for AI-assisted material on human authors and require disclosure of the tool and its purpose; they also identify confidentiality as a constraint on the use of AI in peer review [9]. Policy responses at universities have followed the same concerns without becoming uniform. Dabis and Csáki's analysis of policies from 30 universities identified accountability, human oversight, transparency, and local implementation as recurring themes [10].

The editorial-workflow literature gives a more concrete picture of both potential and limitation. Checco *et al.* developed and evaluated an AI-assisted peer-review approach using 3,300 conference papers and reviews, while noting risks of bias in automated quality judgements [11]. More recently, Thakkar *et al.* ran a randomised study of LLM feedback at ICLR 2025 involving more than 20,000 reviews; reviewers who received feedback produced more informative revisions [12]. Those results concern a specific feedback intervention, not a general replacement for human review. Integrity concerns are equally concrete. Walters and Wilder found fabricated and erroneous entries in bibliographies generated by ChatGPT [13], and Koike *et al.* showed that ordinary task-related prompt constraints can materially change AI-text detector performance [14]. Detection output should therefore not be treated as proof of authorship.

Sentence embeddings provide a complementary tool for comparing abstracts at scale. BGE-M3 supports multilingual and long-document retrieval settings, which matters only if the embedded titles and abstracts are themselves multilingual [6]. The search queries included English, Portuguese, and Spanish terms, but all 711 retained titles and abstracts were English (by language detection). The multilingual setting of BGE-M3 was therefore not needed; the embeddings use only its English encoding. Embedding quality on general benchmarks does not, by itself, establish validity for a specific downstream construct; construct-validity evidence must be gathered for the intended interpretation [7]. For that reason, the semantic dimensions in this paper are theory-guided and their interpretation is limited to the written poles and observed patterns. A blinded two-rater human-validation study was conducted to evaluate their convergent validity (Section IV-D).

III. MATERIALS AND METHODS

A. Corpus construction and scope screening

The starting corpus contained 6,261 articles and reviews with publication years from 2020–2026, assembled from OpenAlex, Europe PMC, Scopus, Web of Science, PubMed/MEDLINE, and BibTeX exports from ACM/IEEE-related searches. Records were assembled in a locally preserved April 2026 snapshot. Local processing artifacts are dated 10 April and manual exports include an early-11-April timestamp; complete execution logs and record-level retrieval timestamps were not retained. The supplement distinguishes preserved query specifications from execution logs. Focused scope screening began on 8 August 2026; final anchor configuration and analysis were refined during August. The corpus is therefore analysed as an April 2026 multi-source snapshot rather than as an exhaustive census. Records were deduplicated by DOI and normalised title, and records without usable abstracts, retractions, and records outside the time window were excluded before screening.

For the article population, we applied documented lexical-first scope rules to title and abstract. Positive signals covered scholarly writing, publishing, editorial and peer-review workflows, research integrity, AI-text detection, evidence-synthesis automation, and AI literacy connected to higher education. Negative signals covered clinical applications, engineering/industrial applications, language learning, and school-level pedagogy. Conflicts were routed through category-aware rules and a recorded review ledger. Cluster membership was retained only as context and never determined inclusion.

The resulting ledger contained 711 included records, 5,542 excluded records, and eight unresolved borderline records. Table II reports the complete processing and decision ledger. The unresolved records were excluded from all analyses reported here. This is single-reviewer, title-and-abstract screening; no inter-rater reliability is claimed for this manuscript. Table I gives the retained corpus profile. The corpus is strongly concentrated in recent publication years: 20 included records predate 2023 and 691 were published in 2023–2026. This distribution may reflect field emergence, retrieval coverage, and database indexing; it is reported without trend modelling.

B. Theory-guided semantic dimensions

We used the 1,024-dimensional dense SentenceTransformer output of BGE-M3, with L2-normalised document and prototype embeddings [6]. The archived document representation joined the title, a period and space, and the first 6,000 characters of the abstract. The stored model configuration permits 8,192 tokens; longer inputs are truncated by the encoder. Each final pole contains five controlled

TABLE I

RETAINED CORPUS AFTER SCOPE SCREENING ($N = 711$), BY PUBLICATION YEAR AND SOURCE. THE 2026 COUNT IS PROVISIONAL.

Year	n	Source	n
2020	4	Web of Science export	533
2021	6	PubMed/MEDLINE export	87
2022	10	Europe PMC	40
2023	99	ACM/IEEE BibTeX export	40
2024	196	Scopus	9
2025	314	OpenAlex	2
2026	82		

Note: 2026 is a partial publication/indexing year and is used for context only. Counts describe the project corpus, not the annual population of all publications. After deduplication by DOI and normalised title, each record was assigned to a single source using the first-seen priority order in which databases were processed (Web of Science, PubMed/MEDLINE, Europe PMC, ACM/IEEE, Scopus, OpenAlex). The distribution therefore reflects both source coverage and this assignment rule.

TABLE II

CORPUS CONSTRUCTION AND FINAL SCOPE DECISIONS.

Step	Input	Removed	Retained
<i>Stage A. Broad-corpus assembly</i>			
Source assembly	6,752 + 1,648	–	8,400
Initial retrieval filters	8,400	–1,821	6,579
Cleaning and embedding eligibility	6,579	–318	6,261
<i>Stage B. Article-specific scope screen</i>			
Title-and-abstract scope screen	6,261	–5,542; –8 pending	711

Note: Initial filters: no abstract 1,241; off-topic title 150; weak alignment 228; low relevance 202. Cleaning: outside period 31; retracted 20; off-topic 45; no usable embedding 222. Stage B is single-reviewer title-and-abstract screening.

abstract prototypes. For axis $A \in \{T, G\}$, let $\mathbf{p}_{A,\pm,j}$ be their normalised embeddings and define

$$\mathbf{c}_{A,\pm} = \frac{1}{5} \sum_{j=1}^5 \mathbf{p}_{A,\pm,j}, \quad (1)$$

$$\mathbf{m}_A = (\mathbf{c}_{A,+} + \mathbf{c}_{A,-})/2, \quad (2)$$

$$\hat{\mathbf{u}}_A = \frac{\mathbf{c}_{A,+} - \mathbf{c}_{A,-}}{\|\mathbf{c}_{A,+} - \mathbf{c}_{A,-}\|_2}, \quad (3)$$

$$s_A(\mathbf{v}) = (\mathbf{v} - \mathbf{m}_A) \cdot \hat{\mathbf{u}}_A. \quad (4)$$

Pole centroids are not re-normalised after averaging. Scores are continuous projections, not probabilities or topic memberships. The supplement records the checkpoint, preprocessing, available runtime information, and the command that applies the exact final anchors to the preserved document embeddings.

Technological specificity (T) ranges from generic AI/ML without a named generative model to papers in which

a named model or model family (e.g., ChatGPT, GPT-4, Claude, Gemini, or LLaMA) is central. **Governance orientation (G)** ranges from workflow support for writing, searching, review, and synthesis to integrity, authorship, disclosure, policy, detection, and institutional safeguards. Figure 1 illustrates the construction and interpretation of a semantic axis using T as the running example.

The 20 prototypes (5 per pole for each axis) anchor the measurement. Because prototype wording can influence the resulting axis, each prototype was written in several candidate variants. The final set was chosen by a greedy search: each of the 20 positions takes the variant that minimises the absolute $T \times G$ axis cosine (the magnitude of the cosine between the two unit vectors), subject to leave-one-prototype-out minimum Spearman $\rho \geq 0.97$ for both axes. The objective uses only the two anchor vectors. The stability constraint uses corpus document-score ranks; neither the document-level T-G association nor the human ratings are an objective or acceptance criterion of the final selector. BGE-M3 embeddings for all candidate prototypes were pre-computed, and the search continues until no position can be improved.

The absolute $T \times G$ axis cosine is 0.0007 at the prototype level. Because the objective minimised this quantity, near-orthogonality is a property of the selection procedure rather than an independent finding. The document-level association is estimated using the frozen final configuration; no coefficient threshold was imposed by the final selector. Leave-one-prototype-out stability is excellent for both axes (T min $\rho = 0.983$, G min $\rho = 0.989$). The supplementary material records the candidate variants, final configuration, and scoring procedure.

A near-zero prototype cosine does not guarantee independent document scores. The Pearson correlation between T and G is $r = 0.162$ in the retained corpus: the dimensions are nearly orthogonal at the prototype level but weakly correlated at the document level.

C. Analysis plan

We report annual counts and score distributions for four completed calendar-year cohorts: 2020–2022 ($n = 20$), 2023 ($n = 99$), 2024 ($n = 196$), and 2025 ($n = 314$). The 2026 records ($n = 82$) are shown for completeness but excluded from comparisons between completed calendar-year cohorts because 2026 is incomplete at the snapshot date. Completed calendar years do not imply exhaustive database coverage. For the principal association, we estimated an ordinary least-squares model of G on T and publication year, with sensitivity analyses excluding 2026 and restricted to 2023–2025 with categorical year adjustment. Coefficient intervals use 2,000 paired-record percentile bootstrap resamples; seeds for the principal and sensitivity models are reported in




Fig. 1. Construction and interpretation of a prototype-based semantic axis, illustrated with the technological-specificity (T) dimension. (a) Five prototype abstracts are written at each pole and embedded with BGE-M3; their centroids define the axis direction. (b) A document's score is its centred projection onto the unit vector between centroids: $s = (\mathbf{v}_{\text{doc}} - \mathbf{m}) \cdot \hat{\mathbf{u}}$. (c) The resulting 1-D score distribution for the 711-record corpus, with leave-one-out prototype stability (minimum $\rho = 0.983$) confirming that the axis is not determined by any single prototype.

the supplementary reproduction manifest (S10). The partial correlation is calculated after residualising both T and G on publication year. These models describe association rather than causal effects.

D. Human validation protocol

A blinded 130-record sample was drawn from the included corpus, stratified by publication-year cohort and the automated-score range available at sampling; 129 records had complete matched ratings on both axes. The sample and human ratings were retained when final automated scores were reconciled. Because the sample was stratified partly by automated-score range, the human-automated correlations reported in Section IV-D are correlations within the stratified validation sample. Population inference would require

appropriate sampling weights and additional assumptions. Two independent raters scored T and G on anchored 1–5 Likert scales without seeing author, venue, year, source, or automatic score. The protocol specifies rank-order inter-rater consistency (Spearman), Spearman association between the automated score and the mean of the two human scores, calibration by score quintile, and bootstrap confidence intervals from 10,000 paired-record resamples; the seed is reported in S10. Final scores were matched to the ratings through the preserved form-to-record crosswalk. Full validation results are reported in Section IV-D.

IV. RESULTS

A. Concentration after 2023

Figure 2(a) shows that the corpus contains four records dated 2020, 314 dated 2025, and 82 dated 2026. The 2026 count is partial and database-dependent. Annual counts are not a complete census because database coverage, search histories, and indexing lag differ by source. The clear post-2023 concentration is therefore reported as a corpus-level pattern, not as an estimate of total annual field output.

The corpus is not limited to higher-education commentary. It includes AI for scientific writing and publishing, editorial workflows, research-integrity policies, AI-text detection, and evidence-synthesis automation. This broader but still delimited definition is important: an AI tool used to screen studies in a health domain is in scope when the research object is the scholarly workflow, whereas an AI system used to predict a clinical outcome is not.


B. Semantic composition

The April 2026 snapshot is concentrated in 2023–2025 (609/711 records). The 2020–2022 baseline contains 20 records and 2026 is partial; cohort summaries therefore describe the retained corpus rather than estimate a stable temporal trend. In the pooled 2020–2022 baseline, median T was -0.067 and median G was -0.111. Across the completed annual cohorts from 2023 to 2025, median T was 0.026 in 2023, 0.007 in 2024, and -0.018 in 2025; median G was 0.081, 0.047, and 0.033, respectively. Figures 2(b) and 2(c) show the overlapping cohort distributions of T and G. These distributions do not establish a smooth linear change over publication years.

C. Named-model specificity and governance orientation

Technological specificity and governance orientation are positively associated in the retained corpus (Figure 3). Controlling for publication year (encoded as a continuous integer), T predicted G (unstandardised $\beta = 0.206$, 95% CI [0.119, 0.300]). The partial correlation was $r = 0.165$ ($p = 1.1 \times 10^{-5}$); the model explained 2.8% of G's variance, compared with 0.07% for year alone. The unique effect of T was small ($f^2 = 0.028$). The association remained positive when 2026 was excluded ($\beta = 0.234$, 95% CI [0.140, 0.324]) and in the complete 2023–2025 cohorts with categorical year adjustment ($\beta = 0.187$, 95% CI [0.095, 0.284]). This cross-sectional association between embedding-derived indicators does not imply that using a named model causes governance concern.

One reading of this pattern is: papers that make ChatGPT or another named model central more often also address disclosure, integrity, institutional policy, or safeguards. Workflow-oriented work remains present across the score



*2026 is partial and descriptive only; boxes show median and IQR.

Fig. 2. Retained records by publication year and T/G score distributions ($N = 711$). Distribution cohorts are 2020–2022 ($n = 20$), 2023 ($n = 99$), 2024 ($n = 196$), 2025 ($n = 314$), and incomplete 2026 ($n = 82$). Boxes show medians and interquartile ranges; points show individual records. The 2026 cohort is descriptive only. Counts describe the April 2026 corpus snapshot, not exhaustive annual field output. Inputs and production dates are recorded in supplementary manifest S10.

range and should not be cast as displaced by governance research.

D. Human-validation results

Two blinded raters independently assessed 130 sampled records; 129 had complete matched ratings on both axes. All automated scores below use the frozen final anchors and the same document scores as the principal analysis.




Fig. 3. Final T and G scores for all 711 retained records. Higher T indicates greater named-model specificity; higher G indicates greater governance orientation. Colours distinguish 2020–2022 ($n = 20$) from 2023–2026 ($n = 691$, including incomplete 2026). The line is the fit from $G \sim T + \text{publication year}$ evaluated at mean publication year. The coefficient interval is a 95% percentile bootstrap interval; no uncertainty band is plotted. Inputs and production dates are in S10.

Inter-rater rank consistency. Spearman's ρ measures rank association rather than exact agreement; quadratic-weighted κ is reported alongside it. For technological specificity (T), inter-rater Spearman $\rho = 0.755$ (95% bootstrap CI [0.659, 0.834], $p < 0.001$), Pearson $r = 0.750$, and quadratic-weighted $\kappa = 0.631$. For governance orientation (G), inter-rater Spearman $\rho = 0.806$ (95% bootstrap CI [0.722, 0.871], $p < 0.001$), Pearson $r = 0.823$, and quadratic-weighted $\kappa = 0.805$.

Convergence with final automated scores. The reference was the mean of the two human ratings. For T, human–automated Spearman $\rho = 0.755$ (95% bootstrap CI [0.663, 0.821], $p < 0.001$; Pearson $r = 0.775$). Per-rater Spearman correlations were 0.607 and 0.783. For G, human–automated Spearman $\rho = 0.623$ (95% bootstrap CI [0.490, 0.732], $p < 0.001$; Pearson $r = 0.644$). Per-rater Spearman correlations were 0.516 and 0.656.

Quintile calibration. Records were grouped into quintiles of the final automated scores within the validation sample. For T, mean human scores across quintiles 1–5 were 1.73, 1.79, 2.76, 3.56, and 4.35. For G, mean human scores across

quintiles 1–5 were 1.42, 3.33, 3.26, 4.19, and 4.38. Human scores rose overall, but the quintile means were not strictly monotonic for both axes.

These unweighted results describe the score-stratified validation sample. They support convergence with the judgments of two human readers; they do not estimate population-wide validity or establish validity for other corpora.

V. DISCUSSION

The weak association between named-model specificity and governance orientation is plausible in scholarly communication: institutions and publishers must address authorship, disclosure, assessment, and reliability as models enter writing and review workflows.

Scope definition is essential. The broad AI search returned thousands of clinical, engineering, and teaching records. The documented screening procedure and preserved internal ledger define the boundary around scholarly workflows and retain unresolved cases outside the analytical corpus.

The semantic axes complement keyword co-occurrence, bibliographic coupling, and co-citation methods by positioning records on theory-guided dimensions. Findings are constrained to the constructs represented by the chosen poles.

Interpretation is constrained by several features of the design. The corpus is not an exhaustive census and is dominated by recent records and Web of Science exports. Abstract-based screening cannot always distinguish a scholarly task from an incidental method reference. Prototype wording and selection affect the measurement. The near-zero axis cosine is an optimisation result; LOO stability concerns robustness to individual anchors. Neither establishes construct validity. Human validation supplies additional evidence, but its unweighted correlations describe the deliberately stratified sample. Corpus coverage and the two-rater design limit generalisation.

The publication supplement provides exact prototype texts, de-identified derived scores, matched validation ratings, and statistical reproduction code. Future work can release additional screening material where source terms permit, compare alternative embedding models, and analyse changes as the 2026 literature matures.

VI. CONCLUSION

This paper maps 711 studies on AI in scholarly communication and research workflows with publication years 2020–2026 in a locally preserved April 2026 snapshot. The 2026 cohort is incomplete and supplies descriptive context; comparisons between completed calendar-year cohorts exclude it. The screened literature is concentrated in 2023–2025 and shows a positive year-adjusted association between named-model specificity and governance orientation. The result is best read as a descriptive map of a developing field

rather than a causal account or a replacement for human judgement. By making scope rules and semantic anchors explicit, the approach supports evidence-based discussion of AI's practical and institutional roles in research.

DECLARATION ON USE OF GENERATIVE ARTIFICIAL INTELLIGENCE

Claude Code using a DeepSeek model assisted code development and data processing during analysis. NotebookLM assisted background research during study preparation. ChatGPT assisted manuscript review, statistical consistency checks, and revision during writing and final preparation. Some candidate prototype texts were drafted or revised with generative-AI tools as measurement inputs. Exact assistance-model versions and candidate-by-candidate generation records were not retained. The authors reviewed the final anchors, data-selection decisions, statistical outputs, cited sources, and manuscript, and take responsibility for the work.

CODE AND DATA AVAILABILITY

The camera-ready paper, source files, and individually browsable reproduction materials are available in the versioned public release. The included derived tables and scripts reproduce the reported statistics and figures. Source database records and document embeddings are omitted under their access terms; recomputing document scores requires those preserved inputs. Historical retrieval timestamps and embedding-runtime versions were not recorded.

REFERENCES

- [1] A. Pritchard, "Statistical bibliography or bibliometrics?," *Journal of Documentation*, vol. 25, no. 4, pp. 348–349, 1969.
- [2] N. Donthu, S. Kumar, D. Mukherjee, N. Pandey, and W. M. Lim, "How to conduct a bibliometric analysis: An overview and guidelines," *Journal of Business Research*, vol. 133, pp. 285–296, 2021, doi: 10.1016/j.jbusres.2021.04.070.
- [3] M. Aria and C. Cuccurullo, "bibliometrix: An R-tool for comprehensive science mapping analysis," *Journal of Informetrics*, vol. 11, no. 4, pp. 959–975, 2017, doi: 10.1016/j.joi.2017.08.007.
- [4] M. J. Page *et al.*, "The PRISMA 2020 statement: an updated guideline for reporting systematic reviews," *BMJ*, vol. 372, p. n71, 2021, doi: 10.1136/bmj.n71.
- [5] M. L. Rethlefsen *et al.*, "PRISMA-S: an extension to the PRISMA statement for reporting literature searches in systematic reviews," *Systematic Reviews*, vol. 10, p. 39, 2021, doi: 10.1186/s13643-020-01542-z.
- [6] J. Chen, S. Xiao, P. Zhang, K. Luo, D. Lian, and Z. Liu, "BGE M3-Embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation," arXiv:2402.03216, 2024.
- [7] Q. Fang, D. Nguyen, and D. L. Oberski, "Evaluating the construct validity of text embeddings with application to survey questions," *EPJ Data Science*, vol. 11, art. no. 39, 2022, doi: 10.1140/epjds/s13688-022-00353-7.
- [8] B. D. Lund and K. T. Naheem, "Can ChatGPT be an author? A study of artificial intelligence authorship policies in top academic journals," *Learned Publishing*, vol. 37, no. 1, pp. 13–21, 2024, doi: 10.1002/leap.1582.
- [9] International Committee of Medical Journal Editors, "Use of artificial intelligence in publishing," *ICMJE Recommendations*, 2026. [Online]. Available: <https://www.icmje.org/recommendations/browse/artificial-intelligence/>. [Accessed: Aug. 8, 2026].
- [10] A. Dabis and C. Csáki, "AI and ethics: Investigating the first policy responses of higher education institutions to the challenge of generative AI," *Humanities and Social Sciences Communications*, vol. 11, art. no. 1006, 2024, doi: 10.1057/s41599-024-03526-z.
- [11] A. Checco, L. Bracciale, P. Loreti, S. Pinfield, and G. Bianchi, "AI-assisted peer review," *Humanities and Social Sciences Communications*, vol. 8, art. no. 25, 2021, doi: 10.1057/s41599-020-00703-8.
- [12] N. Thakkar *et al.*, "A large-scale randomized study of large language model feedback in peer review," *Nature Machine Intelligence*, vol. 8, pp. 326–336, 2026, doi: 10.1038/s42256-026-01188-x.
- [13] W. H. Walters and E. I. Wilder, "Fabrication and errors in the bibliographic citations generated by ChatGPT," *Scientific Reports*, vol. 13, art. no. 14045, 2023, doi: 10.1038/s41598-023-41032-5.
- [14] R. Koike, M. Kaneko, and N. Okazaki, "How you prompt matters! Even task-oriented constraints in instructions affect LLM-generated text detection," in *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024, pp. 14384–14395, doi: 10.18653/v1/2024.findings-emnlp.841.